

# 1. Introduction to Operating Systems

Operating System: Three Easy Pieces

---

Dr. Jianhhui Yue

Slides adapted from Dr. Youjip Won

# What happens when a program runs?

- ▣ A running program executes instructions.
  1. The processor **fetches** an instruction from memory.
  2. **Decode**: Figure out which instruction this is
  3. **Execute**: i.e., add two numbers, access memory, check a condition, jump to function, and so forth.
  4. The processor moves on to the **next instruction** and so on.

# Operating System (OS)

- ▣ Responsible for
  - ◆ Making it easy to **run** programs
  - ◆ Allowing programs to **share** memory
  - ◆ Enabling programs to **interact** with devices

**OS is in charge of making sure the system operates  
correctly and efficiently.**

# Virtualization

- ▣ The OS takes a **physical resource** and transforms it into a **virtual form** of itself.
  - **Physical resource**: Processor, Memory, Disk ...
  - ◆ The virtual form is more general, powerful and easy-to-use.
  - ◆ Sometimes, we refer to the OS as a **virtual machine**.

# System call

- ▣ System call allows user **to tell the OS what to do.**
  - ◆ The OS provides some interface (APIs, standard library).
  - ◆ A typical OS exports a few hundred system calls.
    - Run programs
    - Access memory
    - Access devices

# The OS is a resource manager.

- ▣ The OS **manage resources** such as *CPU*, *memory* and *disk*.
- ▣ The OS allows
  - ◆ Many programs to run → Sharing the CPU
  - ◆ Many programs to *concurrently* access their own instructions and data → Sharing memory
  - ◆ Many programs to access devices → Sharing disks

# Virtualizing the CPU

- ▣ The system has a very large number of virtual CPUs.
  - ◆ Turning a single CPU into a seemingly infinite number of CPUs.
  - ◆ Allowing many programs to seemingly run at once  
→ **Virtualizing the CPU**

# Virtualizing the CPU (Cont.)

```
1      #include <stdio.h>
2      #include <stdlib.h>
3      #include <sys/time.h>
4      #include <assert.h>
5      #include "common.h"
6
7      int
8      main(int argc, char *argv[])
9      {
10         if (argc != 2) {
11             fprintf(stderr, "usage: cpu <string>\n");
12             exit(1);
13         }
14         char *str = argv[1];
15         while (1) {
16             Spin(1); // Repeatedly checks the time and
17                     // returns once it has run for a second
18             printf("%s\n", str);
19         }
20         return 0;
21     }
```

**Simple Example(cpu.c): Code That Loops and Prints**



# Virtualizing the CPU (Cont.)

## ▣ Execution result 1.

```
prompt> gcc -o cpu cpu.c -Wall
prompt> ./cpu "A"
A
A
A
^C
prompt>
```

**Run forever; Only by pressing "Control-c" can we halt the program**

# Virtualizing the CPU (Cont.)

## ▣ Execution result 2.

```
prompt> ./cpu A &  ./cpu B &  ./cpu C &  ./cpu D &  
[1] 7353  
[2] 7354  
[3] 7355  
[4] 7356  
A  
B  
D  
C  
A  
B  
D  
C  
A  
C  
B  
D  
...
```

**Even though we have only **one processor**, all four of  
programs seem to be running **at the same time**!**

# Virtualizing Memory

- ▣ The physical memory is an array of bytes.
- ▣ A program keeps all of its data structures in memory.
  - ◆ **Read memory** (load):
    - Specify an address to be able to access the data
  - ◆ **Write memory** (store):
    - Specify the data to be written to the given address

# Virtualizing Memory (Cont.)

## ■ A program that Accesses Memory (mem.c)

```
1      #include <unistd.h>
2      #include <stdio.h>
3      #include <stdlib.h>
4      #include "common.h"
5
6      int
7      main(int argc, char *argv[])
8      {
9          int *p = malloc(sizeof(int)); // a1: allocate some
                                         memory
10         assert(p != NULL);
11         printf("(d) address of p: %08x\n",
12             getpid(), (unsigned) p); // a2: print out the
                                         address of the memory
13         *p = 0; // a3: put zero into the first slot of the memory
14         while (1) {
15             Spin(1);
16             *p = *p + 1;
17             printf("(d) p: %d\n", getpid(), *p); // a4
18         }
19         return 0;
20     }
```

# Virtualizing Memory (Cont.)

- The output of the program `mem.c`

```
prompt> ./mem
(2134) memory address of p: 00200000
(2134) p: 1
(2134) p: 2
(2134) p: 3
(2134) p: 4
(2134) p: 5
^C
```

- ◆ The newly allocated memory is at address `00200000`.
- ◆ It updates the value and prints out the result.

# Virtualizing Memory (Cont.)

## ▣ Running `mem.c` multiple times

```
prompt> ./mem & ./mem &  
[1] 24113  
[2] 24114  
(24113) memory address of p: 00200000  
(24114) memory address of p: 00200000  
(24113) p: 1  
(24114) p: 1  
(24114) p: 2  
(24113) p: 2  
(24113) p: 3  
(24114) p: 3  
...
```

- ◆ It is as if each running program has its **own private memory**.
  - Each running program has allocated memory at the same address.
  - Each seems to be updating the value at `00200000` independently.

# Virtualizing Memory (Cont.)

- ▣ Each process accesses its own private **virtual address space**.
  - ◆ The OS maps **address space** onto the **physical memory**.
  - ◆ A memory reference within one running program does not affect the address space of other processes.
  - ◆ Physical memory is a shared resource, managed by the OS.

# The problem of Concurrency

- ▣ The OS is juggling **many things at once**, first running one process, then another, and so forth.
- ▣ Modern **multi-threaded programs** also exhibit the concurrency problem.



# Concurrency Example

## ■ A Multi-threaded Program (thread.c)

```
1      #include <stdio.h>
2      #include <stdlib.h>
3      #include "common.h"
4
5      volatile int counter = 0;
6      int loops;
7
8      void *worker(void *arg) {
9          int i;
10         for (i = 0; i < loops; i++) {
11             counter++;
12         }
13         return NULL;
14     }
15
16     int
17     main(int argc, char *argv[])
18     {
19         if (argc != 2) {
20             fprintf(stderr, "usage: threads <value>\n");
21             exit(1);
22         }
```

# Concurrency Example (Cont.)

```
23         loops = atoi(argv[1]);
24         pthread_t p1, p2;
25         printf("Initial value : %d\n", counter);
26
27         Pthread_create(&p1, NULL, worker, NULL);
28         Pthread_create(&p2, NULL, worker, NULL);
29         Pthread_join(p1, NULL);
30         Pthread_join(p2, NULL);
31         printf("Final value : %d\n", counter);
32         return 0;
33     }
```

- ◆ The main program creates **two threads**.
  - Thread: a function running within the same memory space. Each thread start running in a routine called `worker()`.
  - `worker()`: increments a counter

# Concurrency Example (Cont.)

- ▣ `loops` determines how many times each of the two workers will **increment the shared counter** in a loop.

- ◆ `loops: 1000.`

```
prompt> gcc -o thread thread.c -Wall -pthread
prompt> ./thread 1000
Initial value : 0
Final value : 2000
```

- ◆ `loops: 100000.`

```
prompt> ./thread 100000
Initial value : 0
Final value : 143012 // huh??
prompt> ./thread 100000
Initial value : 0
Final value : 137298 // what the??
```

# Why is this happening?

- ▣ Increment a shared counter → take three instructions.
  1. Load the value of the counter from memory into register.
  2. Increment it
  3. Store it back into memory
  
- ▣ These three instructions do not execute **atomically**. → Problem of **concurrency** happen.

# Persistence

- ▣ Devices such as DRAM store values in a volatile.
- ▣ *Hardware* and *software* are needed to store data **persistently**.
  - ◆ **Hardware:** I/O device such as a hard drive, solid-state drives(SSDs)
  - ◆ **Software:**
    - File system manages the disk.
    - File system is responsible for storing any files the user creates.

# Persistence (Cont.)

- Create a file (`/tmp/file`) that contains the string “hello world”

```
1      #include <stdio.h>
2      #include <unistd.h>
3      #include <assert.h>
4      #include <fcntl.h>
5      #include <sys/types.h>
6
7      int
8      main(int argc, char *argv[])
9      {
10         int fd = open("/tmp/file", O_WRONLY | O_CREAT
                        | O_TRUNC, S_IRWXU);
11         assert(fd > -1);
12         int rc = write(fd, "hello world\n", 13);
13         assert(rc == 13);
14         close(fd);
15         return 0;
16     }
```

`open()`, `write()`, and `close()` system calls are routed to the part of OS called the file system, which handles the requests

# Persistence (Cont.)

- ▣ What OS does in order to write to disk?
  - ◆ Figure out **where** on disk this new data will reside
  - ◆ **Issue I/O** requests to the underlying storage device
- ▣ File system handles system crashes during write.
  - ◆ **Journaling** or **copy-on-write**
  - ◆ Carefully ordering writes to disk

# Design Goals

- ▣ Build up **abstraction**
  - ◆ Make the system convenient and easy to use.
  
- ▣ Provide high **performance**
  - ◆ Minimize the overhead of the OS.
  - ◆ OS must strive to provide virtualization without excessive overhead.
  
- ▣ **Protection** between applications
  - ◆ Isolation: Bad behavior of one does not harm other and the OS itself.



# Design Goals (Cont.)

- ▣ High degree of **reliability**
  - ◆ The OS must also run non-stop.
- ▣ Other issues
  - ◆ Energy-efficiency
  - ◆ Security
  - ◆ Mobility

Main frame computer  
M360  
Virtual Machine



Mini computer  
PDP-11  
Unix


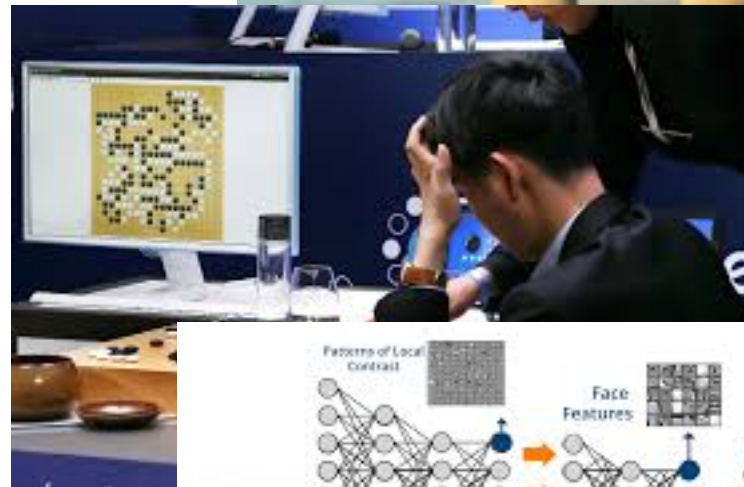
Shell, pipe, signal



- ▣ DOS, Windows, MAC







# What lies all these behind...



# What does OS deal with

## ▣ CPU

- ◆ Execute code with data



## ▣ Memory

- ◆ Read and write code and data



## ▣ Storage

- ◆ Persistently store the code and data





