

4. The Abstraction: The Process

Dr. Jianhui Yue

Slides adapted from Dr. Youjip Won

How to provide the illusion of many CPUs?

▣ CPU virtualizing

- ◆ The OS can promote the illusion that many virtual CPUs exist.
- ◆ **Time sharing**: Running one process, then stopping it and running another
 - By allowing the resource to be used for a little while by one entity, and then a little while by another, and so forth, the resource can be shared by many.
 - The potential cost is **performance**.

▣ Low level mechanism and High level policy

- ◆ Low level methods to implement time sharing
 - Ability to stop running one program and start running another on a given CPU
 - Context switch
- ◆ High level decision
 - Which program should OS run?
 - Scheduling policy

A Process

A process is a **running program.**

- ▣ Machine state
 - ◆ What a program read/write
 - ◆ What parts of the machine affect program execution
- ▣ Comprising of a process:
 - ◆ Memory (address space)
 - Instructions
 - Data section
 - ◆ Registers
 - Program counter(PC)
 - Stack pointer

Process API

- ▣ These APIs are available on any modern OS.
 - ◆ **Create**
 - Create a new process to run a program
 - ◆ **Destroy**
 - Halt a runaway process
 - ◆ **Wait**
 - Wait for a process to stop running
 - ◆ **Miscellaneous Control**
 - Some kind of method to suspend a process and then resume it
 - ◆ **Status**
 - Get some status info about a process


Process Creation

1. **Load** a program code into memory, into the address space of the process.
 - ◆ Programs initially reside on disk in *executable format*.
 - ◆ OS perform the loading process **lazily**.
 - Loading pieces of code or data only as they are needed during program execution.
2. The program's run-time **stack** is allocated.
 - ◆ Use the stack for *local variables*, *function parameters*, and *return address*.
 - ◆ Initialize the stack with arguments → `argc` and the `argv` array of `main()` function

Process Creation (Cont.)

3. The program's **heap** is created.
 - ◆ Used for explicitly requested dynamically allocated data.
 - ◆ Program request such space by calling `malloc()` and free it by calling `free()`.
4. The OS do some other initialization tasks.
 - ◆ input/output (I/O) setup
 - Each process by default has three open file descriptors.
 - Standard input, output and error
5. **Start the program** running at the entry point, namely `main()`.
 - ◆ The OS *transfers control* of the CPU to the newly-created process.

Loading: From Program To Process



Process States

- ▣ A process can be one of three states.

- ◆ **Running**

- A process is running on a processor.

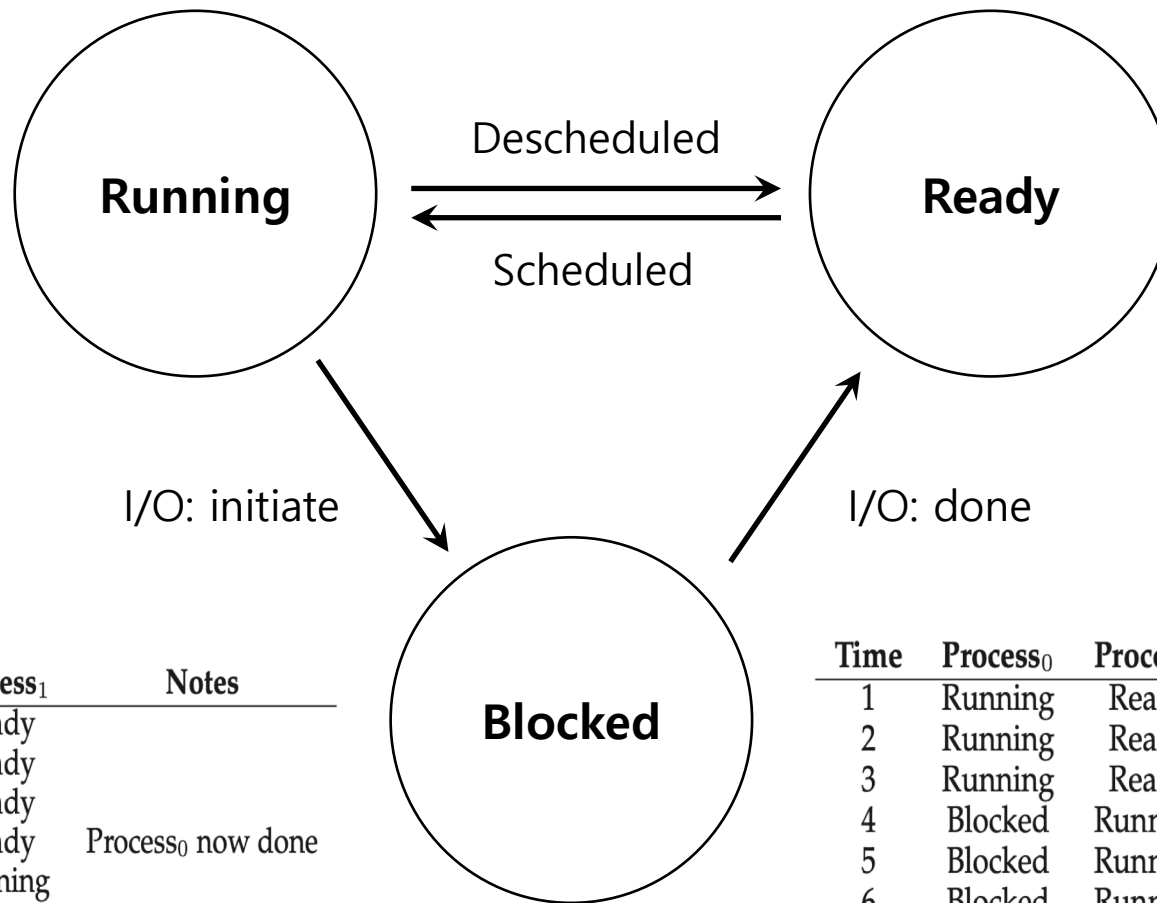
- ◆ **Ready**

- A process is ready to run but for some reason the OS has chosen not to run it at this given moment.

- ◆ **Blocked**

- A process has performed some kind of operation.
 - When a process initiates an I/O request to a disk, it becomes blocked and thus some other process can use the processor.

Process State Transition



Time	Process ₀	Process ₁	Notes
1	Running	Ready	
2	Running	Ready	
3	Running	Ready	
4	Running	Ready	Process ₀ now done
5	–	Running	
6	–	Running	
7	–	Running	
8	–	Running	Process ₁ now done

Figure 4.3: Tracing Process State: CPU Only

Time	Process ₀	Process ₁	Notes
1	Running	Ready	
2	Running	Ready	
3	Running	Ready	Process ₀ initiates I/O
4	Blocked	Running	Process ₀ is blocked, so Process ₁ runs
5	Blocked	Running	
6	Blocked	Running	
7	Ready	Running	I/O done
8	Ready	Running	Process ₁ now done
9	Running	–	
10	Running	–	Process ₀ now done

Figure 4.4: Tracing Process State: CPU and I/O

Data structures

▣ PCB(Process Control Block)

- ◆ A C-structure that contains information **about each process**.
- ◆ **Register context:** a set of registers that define the state of a process
 - a couple of important pieces of information the OS tracks about a process
 - Register context holds the contents of registers for a stopped process
 - Context is stored in memory
 - Resuming a stopped process restore registers from the context

▣ Process list

- ◆ Ready processes
- ◆ Blocked processes
- ◆ Current running processx

Example) The xv6 kernel Proc Structure (Cont.)

```
// the information xv6 tracks about each process
// including its register context and state
struct proc {
    char *mem;           // Start of process memory
    uint sz;             // Size of process memory
    char *kstack;        // Bottom of kernel stack
                        // for this process

    enum proc_state state; // Process state
    int pid;               // Process ID
    struct proc *parent;   // Parent process
    void *chan;            // If non-zero, sleeping on chan
    int killed;            // If non-zero, have been killed
    struct file *ofile[NOFILE]; // Open files
    struct inode *cwd;      // Current directory
    struct context context; // Switch here to run process
    struct trapframe *tf;   // Trap frame for the
                        // current interrupt
};
```

Example) Register Context in xv6

```
// the registers xv6 will save and restore
// to stop and subsequently restart a process
struct context {
    int eip;    // Index pointer register
    int esp;    // Stack pointer register
    int ebx;    // Called the base register
    int ecx;    // Called the counter register
    int edx;    // Called the data register
    int esi;    // Source index register
    int edi;    // Destination index register
    int ebp;    // Stack base pointer register
};

// the different states a process can be in
enum proc_state { UNUSED, EMBRYO, SLEEPING,
                  RUNNABLE, RUNNING, ZOMBIE };
```