

Chapter 14. Memory API

Operating System: Three Easy Pieces

Dr. Jianhui Yue

Slides adapted from Dr. Youjip Won

Memory API: malloc()

```
#include <stdlib.h>

void* malloc(size_t size)
```

- ▣ Allocate a memory region on the heap.
 - ◆ Argument
 - `size_t size` : size of the memory block(in bytes)
 - `size_t` is an unsigned integer type.
 - ◆ Return
 - Success : a void type pointer to the memory block allocated by `malloc`
 - Fail : a null pointer

sizeof()

- ▣ Routines and macros are utilized for `size` in `malloc` instead typing in a number directly.
- ▣ Two types of results of `sizeof` with variables
 - ◆ The actual size of `'x'` is known at run-time.

```
int *x = malloc(10 * sizeof(int));  
printf("%d\n", sizeof(x));
```

4

- ◆ The actual size of `'x'` is known at compile-time.

```
int x[10];  
printf("%d\n", sizeof(x));
```

40


Memory API: free()

```
#include <stdlib.h>

void free(void* ptr)
```


- ▣ Free a memory region allocated by a call to `malloc`.
 - ◆ Argument
 - `void *ptr`: a pointer to a memory block allocated with `malloc`
 - ◆ Return
 - none

Memory Allocating




-----> pointer

```
int *pi; // local variable
```




```
pi = (int *) malloc(sizeof(int) * 4);
```

Memory Freeing



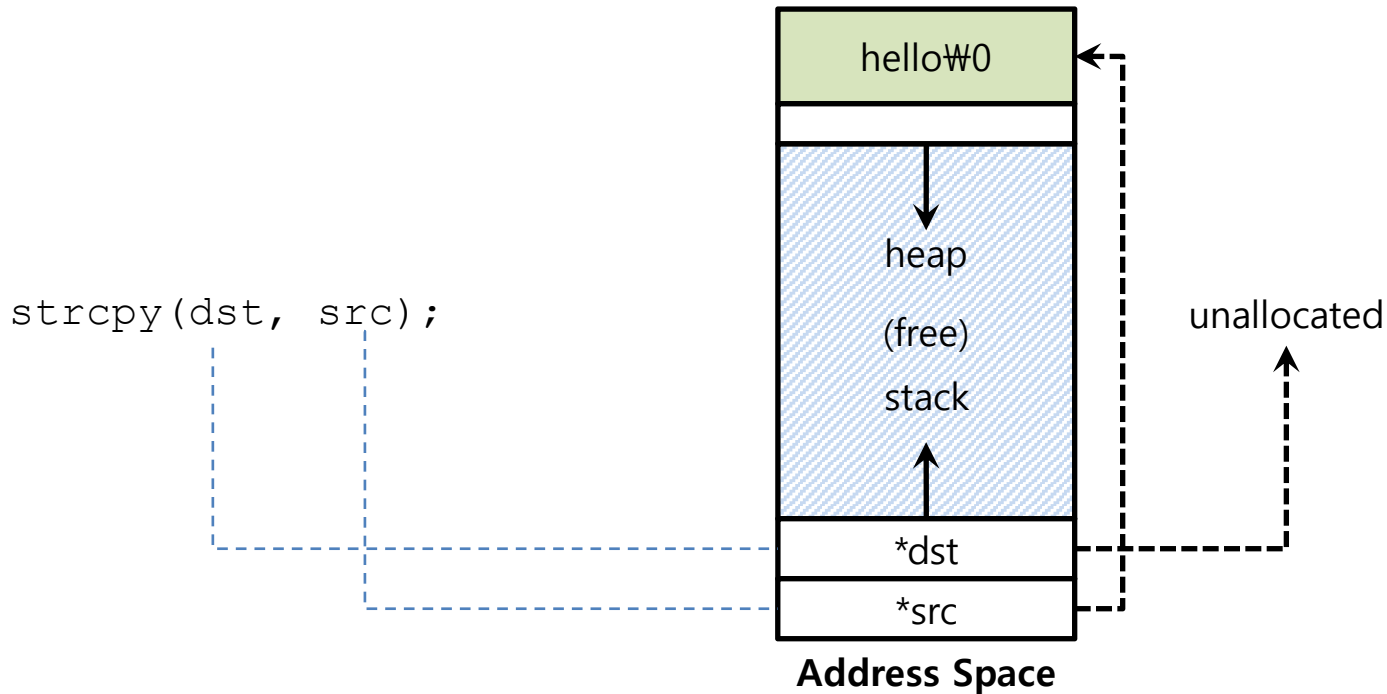
```
free(pi);
```



Forgetting To Allocate Memory

▣ Incorrect code


```
char *src = "hello"; //character string constant
char *dst;           //unallocated
strcpy(dst, src);    //segfault and die
```



Forgetting To Allocate Memory(Cont.)

■ Correct code

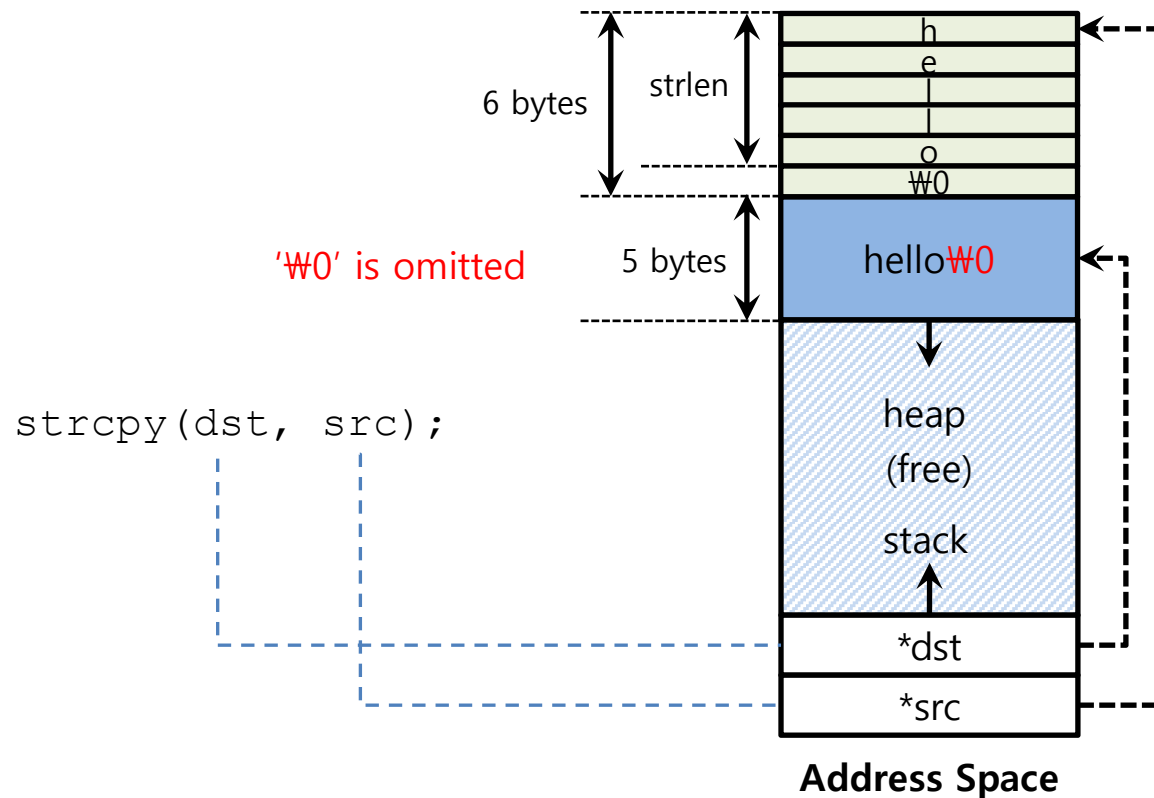
```
char *src = "hello";    //character string constant
char *dst (char *)malloc(strlen(src) + 1 ); // allocated
strcpy(dst, src);        //work properly
```



Not Allocating Enough Memory

❑ Incorrect code, but work properly


```
char *src = "hello"; //character string constant
char *dst (char *)malloc(strlen(src)); // too small
strcpy(dst, src);     //work properly
```




Forgetting to Initialize

■ Encounter an uninitialized read

```
int *x = (int *)malloc(sizeof(int)); // allocated
printf("*x = %d\n", *x); // uninitialized memory access
```



Address Space




Address Space

Memory Leak

- A program runs out of memory and eventually dies.

unused : unused, but not freed



run out of memory

Other Memory APIs: calloc()

```
#include <stdlib.h>


void *calloc(size_t num, size_t size)
```

- Allocate memory on the heap and zeroes it before returning.
 - ◆ Argument
 - `size_t num` : number of blocks to allocate
 - `size_t size` : size of each block(in bytes)
 - ◆ Return
 - Success : a void type pointer to the memory block allocated by `calloc`
 - Fail : a null pointer

Double Free

- Free memory that was freed already.

```
int *x = (int *)malloc(sizeof(int)); // allocated
free(x); // free memory
free(x); // free repeatedly
```



Other Memory APIs: realloc()

```
#include <stdlib.h>

void *realloc(void *ptr, size_t size)
```


■ Change the size of memory block.

- ◆ A pointer returned by `realloc` may be either the same as `ptr` or a new.
- ◆ Argument
 - `void *ptr`: Pointer to memory block allocated with `malloc`, `calloc` or `realloc`
 - `size_t size`: New size for the memory block(in bytes)
- ◆ Return
 - Success: Void type pointer to the memory block
 - Fail : Null pointer

System Calls


```
#include <unistd.h>

int brk(void *addr)
void *sbrk(intptr_t increment);
```



- malloc library call use **brk** system call.
 - ◆ brk is called to **expand** the program's *break*.
 - *break*. The location of **the end of the heap** in address space
 - ◆ sbrk is an additional call similar with brk.
 - Sbrk adjusts the program break value by adding a possibly negative size
 - brk is used to set the break value to the value of a pointer
 - ◆ Programmers **should never directly call** either brk or sbrk.

Sharing memory between two Processes



System Calls(Cont.)

```
#include <sys/mman.h>

void *mmap(void *start, size_t length, int port, int flags,
int fd, off_t offset)
```

- ◆ mmap maps length bytes starting at offset offset from the file specified by the fd into memory, preferably at address start.
- ◆ start address is a hint only, and is usually specified as 0. The actual place where the object is mapped is returned.
- ◆ mmap system call can create **an anonymous** memory region.
 - The mapping is not backed by any file; the fd and offset arguments are ignored.
 - Set flags to be MAP_ANONYMOUS
 - The use of this flag in conjunction with **MAP_SHARED** is only supported on Linux since kernel 2.4.